Paper Review: A Connection Between Score Matching and Denoising Autoencoders

Beomsu Kim*

July 6, 2021

Paper Information.

- Pascal Vincent. A Connection Between Score Matching and Denoising Autoencoders. Neural Computation, 23(7):1661–1674, 2011.

1 Introduction

- $q(x)$: unknown true pdf, $x \in \mathbb{R}^d$.
- $D_n = \{x^{(1)}, \ldots, x^{(n)}\}$: training set of n i.i.d. samples from $q(x)$.
- $q_0(x) = \frac{1}{n}\sum_{i=1}^{n} \delta(\|x - x^{(i)}\|)$: empirical pdf associated with D_n.
- $q_\sigma(\tilde{x} \mid x) = \frac{1}{(2\pi)^{d/2}\sigma^d} e^{-\frac{1}{2\sigma^2}\|\tilde{x}-x\|^2}$: smoothing kernel or noise model.
- $q_\sigma(\tilde{x}, x) = q_\sigma(\tilde{x} \mid x) q_0(x)$: joint pdf.
- $q_\sigma(\tilde{x}) = \frac{1}{n}\sum_{i=1}^{n} q_\sigma(\tilde{x}, x^{(i)})$: Parzen density estimate based on D_n, obtainable by marginalizing $q_\sigma(\tilde{x}, x)$.
- $p(x; \theta)$: density model with parameters θ.
- $J_1 \cong J_2$: means $J_1(\theta)$ and $J_2(\theta)$ have the same set of minimizers.
- $\mathbb{E}_{q(x)}[g(x)] = \int_x q(x) g(x)\, dx$: expectation with respect to distribution $q(x)$.
- $\text{softplus}(x) = \log(1 + e^x)$: will be applied elementwise to vectors.
- $\text{sigmoid}(x) = \frac{1}{1+e^{-x}} = \text{softplus}'(x)$: will be applied elementwise to vectors.
- I : identity matrix.
- $W^\top$: transpose of matrix W.
- W_i : vector for ith row of W.
- $W_{\cdot,j}$: vector for jth column of W.

*Department of Mathematical Sciences, KAIST. Email beomsu.kim@kaist.ac.kr

2 Denoising Autoencoders

- A training input $x \in D_n$ is first corrupted by additive Gaussian noise of covariance $\sigma^2 I$ yielding corrupted input $\tilde{x} = x + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2 I)$. This corresponds to conditional density

$$q_\sigma(\tilde{x} \mid x) = \frac{1}{(2\pi)^{d/2}\sigma^d} e^{-\frac{1}{2\sigma^2}\|\tilde{x}-x\|^2}.$$

- The corrupted version $\tilde{x}$ is encoded into a hidden representation h through an affine mapping followed by a nonlinearity.

$$h = \text{encode}(\tilde{x}) = \text{sigmoid}(W\tilde{x} + b)$$

- The hidden representation h is decoded into reconstruction x^r through affine mapping.

$$x^r = \text{decode}(h) = W^\top h + c$$

- The parameters $\theta = \{W, b, c\}$ are optimized so that the expected squared reconstruction error $\|x^r - x\|^2$ is minimized, i.e., the objective function being minimized by such a denoising autoencoder (DAE) is

$$\begin{aligned} J_{DAE_\sigma}(\theta) &= \mathbb{E}_{q_\sigma(\tilde{x},x)}[\|\,\text{decode}(\text{encode}(\tilde{x})) - x\|^2] \\ &= \mathbb{E}_{q_\sigma(\tilde{x},x)}[\|W^\top \text{sigmoid}(W\tilde{x} + b) + c - x\|^2]. \end{aligned}$$

3 Score Matching

3.1 Explicit Score Matching (ESM)

- Define the energy-based model

$$p(x;\theta) = \frac{1}{Z(\theta)} \exp(-E(x;\theta)), \qquad Z(\theta) = \int_x \exp(-E(x;\theta))\, dx.$$

- Define the score function

$$\Psi(x;\theta) = \frac{\partial \log p(x;\theta)}{\partial x}.$$

- Explicit score matching minimizes

$$J_{ESM_q}(\theta) = \mathbb{E}_{q(x)}\left[\frac{1}{2}\left\|\Psi(x;\theta) - \frac{\partial \log q(x)}{\partial x}\right\|^2\right].$$

3.2 Implicit Score Matching (ISM)

- Define

$$J_{ISM_q}(\theta) = \mathbb{E}_{q(x)}\left[\frac{1}{2}\|\Psi(x;\theta)\|^2 + \sum_{i=1}^{d} \frac{\partial \Psi_i(x;\theta)}{\partial x_i}\right].$$

- Hyvärinen in *Estimation of Non-normalized Statistical Models by Score Matching* shows that

$$J_{ESM_q} \cong J_{ISM_q}.$$

4 Linking Score Matching to the DAE Objective

4.1 Matching the Score of a Non-Parametric Estimator

- Matching $\Psi(x;\theta)$ with the score of Parzen windows density estimator $q_\sigma(\tilde{x})$ yields

$$J_{ESM_\sigma}(\theta) = \mathbb{E}_{q_\sigma(\tilde{x})}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta) - \frac{\partial \log q_\sigma(\tilde{x})}{\partial \tilde{x}}\right\|^2\right].$$

- $q_\sigma(\tilde{x})$ satisfies the regularities conditions for implicit score matching, so

$$J_{ESM_\sigma} \cong J_{ISM_\sigma}.$$

- This equivalence breaks in the limit $\sigma \to 0$, because q_σ no longer satisfies the regularity conditions, and J_{ESM_σ} can no longer be computed (whereas J_{ISM_σ} remains well-behaved).

4.2 Denoising Score Matching (DSM)

- We define the following denoising score matching (DSM) objective

$$J_{DSM_{q_\sigma}}(\theta) = \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta) - \frac{\partial q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\right\|^2\right].$$

- The underlying intuition is that following the gradient Ψ of the log-density at some corrupted point $\tilde{x}$ should ideally move us towards the clean sample x.
- In other words, the model $p(x;\theta)$ should assign higher likelihood to x than $\tilde{x}$.
- Indeed, with the considered Gaussian kernel we have

$$\frac{\partial \log q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}} = \frac{1}{\sigma^2}(x - \tilde{x})$$

  and this direction corresponds to moving from $\tilde{x}$ back towards clean sample x.
- We observe that

$$J_{ESM_\sigma} \cong J_{DSM_{q_\sigma}}.$$

- Also, for an appropriate choice of $p(x;\theta)$,

$$J_{DSM_{q_\sigma}} \cong J_{DAE_\sigma}.$$

Proposition 1.

$$J_{ESM_\sigma} \cong J_{DSM_{q_\sigma}}$$

Proof. Observe that

$$\begin{aligned} J_{ESM_\sigma}(\theta) &= \mathbb{E}_{q_\sigma(\tilde{x})}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta) - \frac{\partial \log q_\sigma(\tilde{x})}{\partial \tilde{x}}\right\|^2\right] \\ &= \mathbb{E}_{q_\sigma(\tilde{x})}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta)\right\|^2\right] - \mathbb{E}_{q_\sigma(\tilde{x})}\left[\left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x})}{\partial \tilde{x}}\right\rangle\right] + C_1 \end{aligned}$$

for a constant C_1 independent of θ. We also have

$$\begin{aligned} \mathbb{E}_{q_\sigma(\tilde{x})}\left[\left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x})}{\partial \tilde{x}}\right\rangle\right] &= \int_{\tilde{x}} q_\sigma(\tilde{x}) \left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x})}{\partial \tilde{x}}\right\rangle d\tilde{x} \\ &= \int_{\tilde{x}} \left\langle \Psi(\tilde{x};\theta), \frac{\partial}{\partial \tilde{x}} q_\sigma(\tilde{x})\right\rangle d\tilde{x} \\ &= \int_{\tilde{x}} \left\langle \Psi(\tilde{x};\theta), \frac{\partial}{\partial \tilde{x}} \int_x q_0(x) q_\sigma(\tilde{x} \mid x)\, dx\right\rangle d\tilde{x} \\ &= \int_{\tilde{x}} \left\langle \Psi(\tilde{x};\theta), \int_x q_0(x) \frac{\partial q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\, dx\right\rangle d\tilde{x} \\ &= \int_{\tilde{x}} \left\langle \Psi(\tilde{x};\theta), \int_x q_0(x) q_\sigma(\tilde{x} \mid x) \frac{\partial \log q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\, dx\right\rangle d\tilde{x} \\ &= \int_{\tilde{x}} \int_x q_0(x) q_\sigma(\tilde{x} \mid x) \left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\right\rangle dx\, d\tilde{x} \\ &= \int_{\tilde{x}} \int_x q_\sigma(\tilde{x}, x) \left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\right\rangle dx\, d\tilde{x} \\ &= \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\right\rangle\right]. \end{aligned}$$

This shows that

$$\begin{aligned} J_{ESM_\sigma}(\theta) &= \mathbb{E}_{q_\sigma(\tilde{x})}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta)\right\|^2\right] - \mathbb{E}_{q_\sigma(\tilde{x})}\left[\left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x})}{\partial \tilde{x}}\right\rangle\right] + C_1 \\ &= \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta)\right\|^2\right] - \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\left\langle \Psi(\tilde{x};\theta), \frac{\partial \log q_\sigma(\tilde{x} \mid x)}{\partial \tilde{x}}\right\rangle\right] + C_1 \\ &= \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta) - \frac{\partial q_\sigma(\tilde{x} \mid x)}{\partial x}\right\|^2\right] - \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\frac{1}{2}\left\|\frac{\partial q_\sigma(\tilde{x} \mid x)}{\partial x}\right\|^2\right] + C_1 \\ &= J_{DSM_{q_\sigma}}(\theta) + C_1 - C_2 \end{aligned}$$

where C_2 is another constant independent of θ. This proves

$$J_{ESM_\sigma} \cong J_{DSM_{q_\sigma}}.$$

□

Proposition 2.

$$J_{DSM_{q_\sigma}} \cong J_{DAE_\sigma}$$

Proof. We choose

$$p(x;\theta) = \frac{1}{Z(\theta)} \exp(-E(x;\theta))$$

where for $\theta = \{W, b, c\}$,

$$E(x;\theta) = -\frac{\langle c, x\rangle - \frac{1}{2}\|x\|^2 + \sum_i \operatorname{softplus}(\langle W_i, x\rangle + b_i)}{\sigma^2}.$$

We then have

$$\begin{aligned}
\Psi_j(x;\theta) &= \frac{\partial \log p(x;\theta)}{\partial x_j} \\
&= -\frac{\partial E(x;\theta)}{\partial x_j} \\
&= \frac{1}{\sigma^2}\left(c_j - x_j + \sum_i \operatorname{softplus}'(\langle W_i, x\rangle + b_i)\frac{\partial(\langle W_i, x\rangle + b_i)}{\partial x_j}\right) \\
&= \frac{1}{\sigma^2}\left(c_j - x_j + \sum_i \operatorname{sigmoid}(\langle W_i, x\rangle + b_i) W_{ij}\right) \\
&= \frac{1}{\sigma^2}(c_j - x_j + \langle W_{\cdot,j}, \operatorname{sigmoid}(Wx + b)\rangle)
\end{aligned}$$

which we can write as the single equation

$$\Psi(x;\theta) = \frac{1}{\sigma^2}(W^\top \operatorname{sigmoid}(Wx + b) + c - x).$$

It follows that

$$\begin{aligned}
J_{DSM_{q_\sigma}}(\theta) &= \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\frac{1}{2}\left\|\Psi(\tilde{x};\theta) - \frac{\partial q_\sigma(\tilde{x} \mid x)}{\partial x}\right\|^2\right] \\
&= \mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\frac{1}{2}\left\|\frac{1}{\sigma^2}(W^\top \operatorname{sigmoid}(Wx + b) + c - x) - \frac{1}{\sigma^2}(x - \tilde{x})\right\|^2\right] \\
&= \frac{1}{2\sigma^4}\mathbb{E}_{q_\sigma(\tilde{x},x)}\left[\|W^\top \operatorname{sigmoid}(W\tilde{x} + b) + c - x\|^2\right] \\
&= \frac{1}{2\sigma^4} J_{DAE_\sigma}(\theta)
\end{aligned}$$

and so

$$J_{DSM_{q_\sigma}} \cong J_{DAE_\sigma}.$$

□